

A Leakage-Aware Benchmark for Ultra-Small Language Models on Structured Speech-Command Routing

[Removed for double-blind review]

¹[Removed for double-blind review]

Abstract. *Small language models ($\leq 3B$ parameters) are appealing for on-device speech interfaces, but superficial format compliance often conceals severe semantic failure. We benchmark five instruction-tuned checkpoints on a transcript-to-command task derived from Fluent Speech Commands (FSC) under an explicit two-operation contract with balanced policy abstention (50% calls, 50% abstentions). We compare official speaker-disjoint evaluation against a leakage-audited phrase-disjoint split. Under the strict contract, a transparent lexical baseline reaches 77.40% exact match, outperforming all zero-shot language models. The largest model (Qwen3.5-2B) attains 100% contract validity solely via degenerate universal abstention (exact call recall 0%), while three checkpoints (Qwen2.5-0.5B, TinyLlama-1.1B, SmolLM2-1.7B) suffer 0% contract validity despite emitting valid JSON, and Qwen3.5-0.8B attains 60.15%. We demonstrate that this structural failure is eliminated by in-distribution supervised parameter-efficient adaptation (LoRA), and establish guidelines for evaluating SLMs on structured routing tasks.*

1. Introduction and Scope

Deploying natural language command routing locally on edge hardware requires ultra-small language models (SLMs, $\leq 3B$ parameters). However, evaluating these models zero-shot introduces two severe pitfalls: (i) template leakage across speakers in benchmark splits, and (ii) degenerate policy compliance, where models achieve seemingly high accuracy by trivial abstention.

We evaluate five checkpoints: Qwen2.5-0.5B, Qwen3.5-0.8B, TinyLlama-1.1B, SmolLM2-1.7B, and Qwen3.5-2B. The input is a transcription from Fluent Speech Commands (FSC) [1], and the output is a validated JSON object conforming to a strict two-tool contract: `media_control` and `volume_adjust`, with all other 27 native intents mapped to a derived abstention target. The evaluation is balanced 1:1 (7,478 calls and 7,478 abstentions per split). We evaluate on both FSC’s official speaker-disjoint split ($n = 1,934$) and an audited phrase-disjoint split ($n = 2,296$, 38 template clusters) where no transcription template crosses partitions.

2. Experimental Results

As shown in Table 1 and Figure 1, the lexical control dominates all zero-shot models. Exact Call Recall (R_{call}^{exact}) denotes the ratio of correct calls with exact tool and arguments over gold call instances. Qwen3.5-2B attains 100% validity and 50.00% exact match, but inspection reveals complete policy collapse: it outputs `abstain` on 100% of inputs, yielding $R_{call}^{exact} = 0\%$. Qwen3.5-0.8B is the only model attempting active

Table 1. Phrase-disjoint vs. Official test results (%).

System	Phrase-Disjoint (Unseen Templates)			Official (Speaker Split)		
	Valid	Exact	Exact Call Recall	Valid	Exact	Template Exact
Always abstain	100.00	50.00	0.00	100.00	50.00	75.71
Lexical control	100.00	77.40	54.79	100.00	77.92	89.07
Qwen2.5-0.5B	0.00	0.00	0.00	0.00	0.00	0.00
Qwen3.5-0.8B	60.15	43.60	33.45	73.47	54.65	56.28
TinyLlama-1.1B	0.00	0.00	0.00	0.00	0.00	0.00
SmolLM2-1.7B	0.00	0.00	0.00	0.78	0.00	0.00
Qwen3.5-2B	100.00	50.00	0.00	100.00	50.98	76.11

routing, but exhibits an observed split gap of 11.05 pp from official (54.65%) to phrase-disjoint (43.60%). Because the phrase-disjoint cluster bootstrap 95% CI is wide ([26.02%, 63.35%], over 38 clusters with 9 call and 29 abstain clusters), this gap is descriptive rather than a formal causal test.





Figure 1. Trade-off space on phrase-disjoint test. Marker area $\propto$ exact call recall. The lexical baseline dominates zero-shot SLMs; Qwen3.5-2B collapses to 100% abstentions.

Figure 2 reveals the exact point of structural failure: SmolLM2-1.7B produces 96.47% syntactically parseable JSON, but 0% valid contract objects due to schema deviations (missing action/tool keys or invalid parameters). Unconstrained zero-shot decoding is insufficient for sub-3B models to adhere to rigid task schemas.

3. Discussion, Adaptation, and Conclusions

Two orthogonal techniques address this bottleneck: (i) **Grammar-Constrained Decoding** (e.g., GBNF in `llama.cpp`, Outlines [6]), which forces syntactic compliance by



Stages: 1 Invalid JSON -> 2 Valid JSON but invalid contract -> 3 Valid contract but wrong decision -> 4 Valid contract and exact match. Values shown for slices $\geq 8\%$.

Figure 2. Structural cascade decomposing outputs into: (1) invalid JSON, (2) valid JSON but invalid schema, (3) valid contract but wrong call, and (4) exact match.

construction (collapsing stages 1–2 of Figure 2 to zero, though semantic discrimination still relies on the model); and (ii) **Supervised LoRA Adaptation** [5]: an in-distribution Qwen2.5-0.5B LoRA adaptor reaches 100% contract validity and 100% exact match, demonstrating that the contract is readily learnable with parameter-efficient fine-tuning.

In conclusion, zero-shot structured routing without grammar guidance is fragile in the ultra-small regime. Benchmarks must decouple syntactic validity from semantic utility, audit template leakage, and incorporate transparent non-neural baselines. Extending this evaluation to multi-domain SLU benchmarks like SLURP [2] and compositional parsing corpora like STOP [3] represents an essential next step.

References

- [1] L. Lugosch et al. Speech model pre-training for spoken language understanding. *Interspeech*, 2019.
- [2] E. Bastianelli et al. SLURP: A Spoken Language Understanding Resource Package. *EMNLP*, 2020.
- [3] P. Tomasello et al. STOP: A dataset for spoken task-oriented semantic parsing. *IEEE SLT*, 2023.
- [4] Z. Liu et al. MobileLLM: Optimizing sub-billion parameter language models for on-device use cases. *ICML*, 2024.
- [5] E. J. Hu et al. LoRA: Low-rank adaptation of large language models. *ICLR*, 2022.
- [6] B. T. Willard and R. Louf. Efficient guided generation for large language models. *arXiv:2307.09702*, 2023.